

Weekly Meeting Log

The weekly entries below are aligned with the development stages documented in the final report.

Group	The Optimizers		
Members	Vidhi Marathe (28902)	Anmol Gupta (27257)	Nunavath Saidulu (26534)
Week 1 Tasks: who did what		Completed work	
Vidhi Marathe		Inspected the anonymised token sequences, class counts and length distributions; summarised the A/B imbalance and initial EDA findings.	
Anmol Gupta		Set up the leakage-safe stratified evaluation protocol and built the first TF-IDF n-gram baseline for token sequences.	
Nunavath Saidulu		Compared initial classifier behaviour and documented why token-frequency features alone were insufficient; identified the need for sequence-structure features.	
Next step		Move from EDA and baseline TF-IDF to richer sequence features.	

Week 2 Tasks: who did what		Completed work	
Vidhi Marathe		Implemented deterministic sequence statistics covering length, entropy, repetition, compression and related signals.	
Anmol Gupta		Added static sequence-geometry features and character/word TF-IDF representations, including the SVD feature block.	
Nunavath Saidulu		Developed and tested alternative base learners and organised fold-wise feature preparation so label-dependent quantities stayed inside the training fold.	
Next step		Add positional language-model signals and compare heterogeneous learners.	

Week 3 Tasks: who did what		Completed work	
Vidhi Marathe		Built the positional class-conditional n-gram language model with smoothing, position bins and likelihood-ratio summaries.	
Anmol Gupta		Integrated the heterogeneous base learners into a stacking pipeline and generated the out-of-fold matrix for the LogisticRegression meta-learner.	
Nunavath Saidulu		Ran the controlled ablation on raw Markov scores, investigated the scale problem, and recorded the before/after results for the report.	
Next step		Complete stacking, run the controlled ablation, and prepare final evaluation.	

Week 4 Tasks: who did what		Completed work	
Vidhi Marathe		Re-ran the full stack over five random seeds and averaged the test probabilities to produce the final bagged prediction file.	
Anmol Gupta		Consolidated the submission history, prepared the confusion-matrix/error-analysis material, and wrote the Results and Conclusion sections.	
Nunavath Saidulu		Reviewed the final pipeline, submission metrics and limitations; consolidated the final report and submission package.	
Next step		Validate the final submission, consolidate results, and finish the report.	